# Asim Adnan Eijaz


Sydney, NSW
https://asimadnan.com
Email : asimadnaneijaz@hotmail.com
Mobile : 0411 802 251


## Experience

- **Stockland** — Sydney, NSW
  *Senior Artificial Intelligence & Machine Learning Engineer* — *Dec 2022 – Present*
  - Partner with senior stakeholders across research, acquisitions, finance, and marketing to scope high-impact LLM use cases, rapidly prototype solutions, and guide teams from initial concept through to production deployment.
  - Design and deliver end-to-end LLM-powered applications that transform unstructured legal, tax, and commercial documents into validated structured datasets, incorporating human-in-the-loop workflows, citation grounding, and auditability for enterprise use.
  - Led development of Terra Draw, working closely with domain experts to translate complex CAD-based planning workflows into an automated GenAI system that extracts geospatial data and integrates directly into downstream pricing and sales systems.
  - Built and iterated on agent-based site intelligence workflows, collaborating with acquisition teams to define requirements and delivering multi-agent systems that synthesize internal and external data into actionable feasibility insights.
  - Improved trust and adoption of ML systems by introducing explainability layers combining feature attribution and LLM-generated reasoning, enabling non-technical stakeholders to interpret and act on model outputs.
  - Developed executive-facing AI intelligence workflows that surface risk signals and operational insights across 80+ assets, delivering concise, decision-ready summaries for leadership teams.
  - Partnered with business teams to deploy Voice-of-Customer pipelines, using LLM-based classification and sentiment analysis to convert large-scale unstructured feedback into structured insights used across multiple functions.
  - Delivered reusable document intelligence patterns using agent-based extraction and validation, reducing manual analysis effort and enabling faster decision-making in capital transaction workflows.
  - Established evaluation and reliability frameworks for LLM applications, including test set design, failure analysis, and prompt iteration, ensuring production systems meet accuracy and trust requirements.
  - Lead technical workshops, demos, and executive briefings to educate stakeholders on LLM capabilities, best practices, and deployment strategies, translating complex systems into clear business value.
- **Wallace** — Melbourne, VIC
  *Machine Learning Engineer* — *Feb 2022 – Jun 2022*
  - Fine-tuned transformer-based audio models under varying noise conditions, reducing false activations by 30%.
  - Built a monitoring and experiment-repro framework with TensorBoard and config management. Cut debugging and reproduction time by 50%.
  - Packaged wake-word detection into a production-ready Docker library with CI/CD.
- **Digital Health CRC** — Sydney, NSW
  *Data Science Intern* — *Aug 2020 – Dec 2020*
  - Used ML techniques to model opioid adherence and overdose risk. Research published in PLOS ONE.
  - Estimated that targeted interventions for the highest-risk 10% could reduce overdose events by 10.4%.
- **Tkxel** — Lahore, Pakistan
  *Software Engineer* — *Jun 2016 – Dec 2018*
  - Designed and implemented distributed backend architectures and microservices, improving system reliability by 80% and reducing cycle time by 40%.
  - Mentored junior engineers and guided adoption of TDD, CI/CD, and API design best practices.

## Education

- **Macquarie University** — Sydney, NSW
  *Masters in Data Science & Masters in Research* — *Jan 2019 – Dec 2021*
- **GIK Institute of Engineering Sciences and Technology** — Swabi, Pakistan
  *Bachelors in Computer Science* — *Aug 2012 – May 2016*

## Selected Projects

- **Enterprise Content Enrichment Platform**: Designed multi-agent orchestration using Semantic Kernel and Model Context Protocol (MCP) to convert unstructured documents into structured metadata and knowledge graphs. Built ingestion, validation, and retrieval pipelines enabling downstream discovery and search applications.
- **LLM Evaluation Toolkit**: Developed evaluation pipelines for GenAI systems using LLM judges, heuristic metrics, and LangFuse telemetry to benchmark output quality, detect hallucinations, and iterate on prompts and retrieval strategies across multiple use cases.
- **Speech-to-Insight Pipeline**: Built speech-to-text and LLM pipelines that extract topics, sentiment, and action items from sales conversations. Enabled structured insight generation from call recordings to support coaching and performance analytics.
- **GenAI Deployment Patterns (GitHub Projects)**: Maintain public repositories demonstrating production-ready GenAI workflows including document extraction, agent orchestration, retrieval pipelines, and structured data generation. Focus on reproducible architectures that translate LLM capabilities into real business workflows.
- **Supervised Fine-Tuning of Phi for Urdu Tasks**: Experimented with supervised fine-tuning of the Phi model on Urdu-language instructions using a curated dataset. Evaluated improvements in structured extraction and instruction following for low-resource language tasks.
- **Opioid Adherence Modeling**: Published ML research in PLOS ONE on adherence and overdose risk modeling using predictive machine learning methods.

## Technical Skills

- **GenAI & ML**: LLM applications, retrieval systems, RAG patterns, prompt architecture, evaluation pipelines, safety checks, supervised/unsupervised learning, deep learning
  **Architectures**: Cloud-native design, event-driven systems, distributed services, API design, observability
  **Languages**: Python, TypeScript, Ruby, R, SQL, C++
  **Platforms**: AWS (Bedrock, SageMaker, Lambda, ECS), Semantic Kernel, LangChain, MCP, Docker, GitHub Actions, Hugging Face
  **Other**: Rapid prototyping, stakeholder facilitation, technical discovery, enablement